

RavenStack Growth and Retention Memo

RavenStack — Growth & Retention Review and Recommendation

To: Product Leadership **From:** Abi (Growth/Product) **Date:** January 2025 **Re:** Where retention is breaking, and what I'd do about it

Executive summary

We churn 22% of accounts (110 of 500 over two years). The churn is concentrated, not general: DevTools accounts churn at 31%, event-sourced accounts at 30.2% (vs 14.6% for partner-sourced), and mid-size teams (6–20 seats) at 25%. Two things I expected to explain this don't: churned accounts use the product about as much as retained ones (no pre-churn usage drop, equal CSAT, similar ticket volume), and no individual feature separates churners from stayers. What customers tell us is feature gaps (19% of churn events, the #1 code, and the clear leader for Enterprise), followed by budget and support. Revenue-wise the exposure is lopsided: roughly \$14.2M ARR sits on churned subscriptions and **Enterprise is 78.6% of it (\$11.12M)** even though Enterprise churn *rate* is unremarkable.

I recommend two things: (1) a feature-fit discovery loop for Enterprise (interviews + a structured churn-exit survey, feeding a roadmap decision — the data can't tell us *which* capability is missing, only that fit, not usage, is the stated problem), and (2) an experiment on event-sourced accounts targeting the month 4–6 window where their retention curve starts sliding versus partner. Neither is a dashboard. Both are cheap relative to \$14M of recurring revenue at risk.

1. Problem

Leadership asked: where is retention breaking down, what behaviours are associated with churn, which segments deserve attention, and what should we do. Scope: all 500 accounts, Jan 2023 – Dec 2024, ~5,000 subscriptions, 25,000 usage events, 2,000 support tickets, 600 churn events. Metric definitions are pinned in `analysis/README.md` — worth skimming, because this dataset has three different things called “churn” and they disagree (22.0% account churn vs 9.7% subscription churn vs 600 event rows).

2. Key findings

Retention breaks along acquisition and vertical, not plan.

segment	churn	n
DevTools	31.0%	113
Event-sourced	30.2%	96
Partner-sourced	14.6%	89
Cybersecurity	16.0%	100
6–20 seats	25.0%	188
Plan tiers (Basic/Pro/Ent)	21.9–22.1%	—

The event-vs-partner gap is statistically real (chi-sq 5.55, $p=0.018$); the omnibus tests across all industries/sources are borderline ($p=0.066/0.079$), which at $n=500$ I read as “strong signal, modest sample,” not noise. The intersection is ugly: DevTools × event churns at 43.5% ($n=23$, directional only).

It is not an engagement problem. Churned accounts average 21.2 usage events vs 23.4 for retained, touch 14.9 vs 16.0 distinct features, and their usage is flat right up to the churn date (avg usage/event count 9.9 in the last 30 days vs 10.0 in days 31–90 before). CSAT is 4.0 vs 3.95. First response is actually *faster* for churned accounts. Support escalations are the one modest mover (20.9% of churned accounts hit by ≥ 1 escalation vs 17.4%). I went looking for an activation signal in first-30-day feature breadth and found nothing: 24.0% churn at 0–2 features vs 25.2% at 3–5. No cliff, no gradient, no cutoff I'd defend.

Customers say it's fit and money. Reason codes on 600 churn events: features 19.0%, budget 17.3%, support 17.3%, unknown 15.8%, competitor 15.3%, pricing 15.2%. For Enterprise, features is the #1 code by a wider margin; Basic and Pro lead with support. Budget + pricing together are a third of all events.

Event-sourced accounts decay late, not early. Their pooled activity retention matches partner-sourced accounts through ~month 5 (~85–89%) and slides from month 6 on, ending near 75% at M9–M11 while partner holds ~90%. If onboarding (day 0–30) were the problem, the curves would split early. They don't.

3. Customer and revenue impact

\$1.18M MRR / \$14.15M ARR sits on churned subscriptions. Enterprise accounts for \$11.12M of that, Pro \$2.16M, Basic \$0.87M — churn *rates* are near-identical across plans, so the difference is purely deal size. DevTools carries the most at-risk ARR by industry (\$8.0M). Median tenure at churn is 151 days and two-thirds of churners leave inside 180 days, so whatever we fix has to act inside the first two quarters of the relationship.

4. Root-cause hypotheses (hypotheses, not findings)

- **H1 — expectation mismatch at acquisition.** Event-sourced accounts were sold a story at a booth/webinar that the product doesn't operationalise for them, and value erodes once the initial use case stalls (consistent with the late slide, flat usage — they don't disengage, they just don't renew).
- **H2 — feature-fit gap for power/Enterprise use cases.** Consistent with the #1 reason code and Enterprise revenue concentration. The data cannot say *which* capability is missing; feedback text is canned (“missing features”, “switched to competitor”).
- **H3 — mid-market packaging/price friction.** Budget+pricing $\approx$ 32% of events and the 6–20 seat band churns most; consistent with seats-based pricing biting teams as they grow.

5. Recommendation

Two tracks, deliberately not “build feature X”:

1. **Feature-fit discovery loop (Enterprise, 4 weeks).** Structured churn-exit interviews for every Enterprise churn going forward (we get ~34/yr), a 3-question exit survey replacing the current free-text field, and a synthesis deck answering: are the missing capabilities one roadmap theme or scattered? Decision gate: if 2+ of the next 5 Enterprise churn events cite the same theme, it goes to roadmap planning with the \$11M exposure attached to it.
2. **Event-source lifecycle experiment (below).** Cheaper, faster, and it targets the largest churn-rate gap we can actually manipulate.

Explicitly deprioritized: usage-volume nudges (“re-engagement” campaigns) — the data argues churners look engaged; onboarding-day-one rework — the divergence starts around month 4–6, not month 1; and a churn-prediction model — we don't have a prediction problem - we have a decision problem.

6. Proposed experiment

Hypothesis: event-sourced accounts that receive a structured value checkpoint in months 4–6 (usage review against their stated goals, captured at signup, plus a live “are we delivering” 20-min call offer) will retain materially better at M9 than event-sourced accounts without it.

- **Who:** new event-sourced signups (all industries), randomised at signup, 50/50.
- **Control:** current journey. **Treatment:** month-4 value checkpoint + month-6 check-in.
- **Primary metric:** M9 activity retention (currently ~75% for event-sourced vs ~90% partner — the gap we’re buying back).
- **Secondary:** MRR retention on the cohort; reason-code mix at churn (does “features”/“budget” share drop?).
- **Guardrails:** support ticket volume per account, CSAT, trial→paid conversion (must not regress).
- **Duration:** 9 months minimum for the primary read (that’s the honest cost of an M9 endpoint). Full power (~400/arm to detect 8pts at $\alpha=0.05$, power 0.8) is unrealistic at our event signup volume (~50/yr), so I’d pre-register a supplementary 90-day proxy: month 4–6 activity drop-off rate, and accept a directional readout on M9. Better to know the limitation now than pretend.

If H1 is wrong and checkpoints don’t move retention, that’s still information — it points spend at packaging (H3) instead.

7. Success metrics (for the overall initiative, not just the experiment)

- Account churn, event-sourced segment: 30% → $\leq 22\%$ (parity with base) within 2 quarters of full rollout
- Enterprise ARR on churned subs: run-rate down 30% within 12 months (via theme identification + roadmap response)
- M9 retention gap, event vs partner: 14pts → ≤ 7 pts
- Zero regression on guardrails (CSAT, trial→paid, ticket volume)

8. Risks and limitations

Synthetic dataset (Rivalytics/RavenStack) — methodology transfer is the point, not the absolute numbers. $n=500$ means segment intersections are directional at best. Correlation $\neq$ causation throughout: acquisition source “associates with” churn, we have not established mechanism. No interview data exists yet — H1–H3 are inferences from behaviour + reason codes. Recent cohorts are right-censored (cohort analysis uses 2023 signups only). 53% of raw usage events predate the account’s signup date (generator artifact); every usage metric here bounds activity to signup → churn, which is the only defensible window. And the dataset’s trial funnel is degenerate (every trial-first account converts — 86/86), so I don’t quote a trial→paid number; upgrade rate (61% of accounts upgrade at least once) is the honest growth metric available.

9. Next steps

1. Approve exit-survey/interview instrument (draft by Friday)
2. Stand up experiment randomisation + tracking before next event campaign pulls signups
3. Re-run this analysis quarterly; the cohort matrix and segment table are already automated in `analysis/`